import pandas as pd
from openpyxl.workbook import Workbook

#reading excel file

In [2]:

#Top 5 rows of the Excel file
df.head()

Out[2]

t[2]:		Date	Ottawa	Toronto West/Ouest	Toronto East/Est	Windsor	London	Peterborough	St. Catharine's	Sudbury	Sault Saint Marie	Thunder Bay	N
	0	1990- 01-03	55.9	49.1	48.7	45.2	50.1	0.0	0.0	56.4	54.8	56.6	
	1	1990- 01-10	55.9	47.7	46.8	49.7	47.6	0.0	0.0	56.4	54.9	56.8	
	2	1990- 01-17	55.9	53.2	53.2	49.6	53.7	0.0	0.0	55.8	54.9	56.8	
	3	1990- 01-24	55.9	53.2	53.5	49.0	52.1	0.0	0.0	55.7	54.9	56.8	
	4	1990- 01-31	55.9	51.9	52.6	48.6	49.1	0.0	0.0	55.6	54.8	56.8	

In [3]:

 $\# Datatype \ and \ Null \ Information \ about \ the \ columns \ in \ the \ Excel \ file \ df.info()$

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9962 entries, 0 to 9961
Data columns (total 20 columns):

Data	columns (total 20 columns):		
#	Column	Non-Null Count	Dtype
0	Date	9962 non-null	datetime64[ns]
1	Ottawa	9962 non-null	float64
2	Toronto West/Ouest	9962 non-null	float64
3	Toronto East/Est	9962 non-null	float64
4	Windsor	9962 non-null	float64
5	London	9962 non-null	float64
6	Peterborough	9962 non-null	float64
7	St. Catharine's	9962 non-null	float64
8	Sudbury	9962 non-null	float64
9	Sault Saint Marie	9962 non-null	float64
10	Thunder Bay	9962 non-null	float64
11	North Bay	9962 non-null	float64
12	Timmins	9962 non-null	float64
13	Kenora	9962 non-null	float64
14	Parry Sound	9962 non-null	float64
15	Ontario Average/Moyenne provinciale	9962 non-null	float64
16	Southern Average/Moyenne du sud de l'Ontario	9962 non-null	float64
17	Northern Average/Moyenne du nord de l'Ontario	9962 non-null	float64

```
In [4]:
        #Creating a new dataframe
        df2 = pd.DataFrame()
In [5]:
        #Choosing rows from the original dataframe with Fuel Type "Diesel" and pasting it into new
        df2 = df.loc[df['Fuel Type']=='Diesel']
In [6]:
        #Datatype and Null Information about the columns in the new dataframe
        df2.info()
       <class 'pandas.core.frame.DataFrame'>
       Int64Index: 1675 entries, 1675 to 3349
       Data columns (total 20 columns):
          Column
                                                         Non-Null Count Dtype
       --- ----
                                                         _____
                                                         1675 non-null datetime64[ns]
           Date
        1
           Ottawa
                                                         1675 non-null float64
        2 Toronto West/Ouest
                                                         1675 non-null float64
        3 Toronto East/Est
                                                         1675 non-null float64
           Windsor
                                                         1675 non-null float64
        4
        5 London
                                                         1675 non-null float64
        6 Peterborough
                                                         1675 non-null float64
        7
                                                         1675 non-null float64
          St. Catharine's
        8
          Sudbury
                                                         1675 non-null float64
        9 Sault Saint Marie
                                                         1675 non-null float64
        10 Thunder Bay
                                                         1675 non-null float64
        11 North Bay
                                                         1675 non-null float64
        12 Timmins
                                                         1675 non-null float64
        13 Kenora
                                                         1675 non-null float64
        14 Parry Sound
                                                        1675 non-null float64
        15 Ontario Average/Moyenne provinciale
                                                        1675 non-null float64
        16 Southern Average/Moyenne du sud de l'Ontario 1675 non-null float64
        17 Northern Average/Moyenne du nord de l'Ontario 1675 non-null float64
                                                         1675 non-null object
        18 Fuel Type
        19 Type de carburant
                                                         1675 non-null object
       dtypes: datetime64[ns](1), float64(17), object(2)
       memory usage: 274.8+ KB
In [7]:
        df2.head()
Out[7]
```

9962 non-null object

9962 non-null

18 Fuel Type

19 Type de carburant

memory usage: 1.5+ MB

dtypes: datetime64[ns](1), float64(17), object(2)

t[7]:		Date	Ottawa	Toronto West/Ouest	Toronto East/Est	Windsor	London	Peterborough	St. Catharine's	Sudbury	Sault Saint Marie	Thunder Bay
	1675	1990- 01-03	49.3	47.6	48.3	46.5	47.2	0.0	0.0	45.4	45.8	46.6
	1676	1990- 01-10	49.5	47.9	48.6	47.1	47.4	0.0	0.0	45.8	46.1	46.6
	1677	1990- 01-17	49.5	48.6	48.6	47.3	47.4	0.0	0.0	47.2	46.1	46.6
	1678	1990- 01-24	50.4	47.9	48.7	47.6	47.7	0.0	0.0	47.2	46.2	47.2

		Date O	ttawa	Toronto West/Ouest	Toronto East/Est	Windsor	London	Peterborough	St. Catharine's	Sudbury	Sault Saint Marie	Thunder Bay
	1679	1990- 01-31	50.4	47.7	48.7	47.6	47.7	0.0	0.0	47.2	46.5	47.3
[8]:		ating a = pd.Da		dataframe me()								
[9]:	df3.	insert(0,"Da	el' price te",df2['l	Date'], T	rue)		y and pasti.	ng it into	new dat	aframe)
)]:	df3.	head()										
)]:		Da	te Tin	nmins								
	1675	1990-01-0	03	47.2								
	1676	1990-01-	10	47.4								
	1677	1990-01-	17	47.7								
	1678	1990-01-2	24	47.7								
	1679	1990-01-3	31	48.4								
	"-						. ,					
]:	#thi #dai df3. df3	s incon ly valu set_ind = df3.r	istendes. ex('Daesampi		rting th	e date d	column i	he week sta. nto daily v				
	#thi #dai df3. df3	s incon ly valu set_ind = df3.r t 10 va tail(10	istendes. ex('Daesampi	cy, convergete, inpide ('D').fi	rting th	e date d	column i					
	#thi #dai df3. df3 #Las df3.	s incon ly valu set_ind = df3.r t 10 va tail(10	istendes. ex('Daesamp. lues.)	cy, convergete, inpide ('D').fi	rting th	e date d	column i					
	#thi #dai df3. df3 #Las df3.	s incon ly valu set_ind = df3.r t 10 va tail(10	istendes. ex('Daesamp. lues.) ate Ti -22	ate', inpile('D').fi	rting th	e date d	column i					
	#thi #dai df3. df3 #Las df3.	s incon ly valu set_ind = df3.r t 10 va tail(10	es. ex('Daesamp. lues.) ate Ti -22 -23	cy, converge to the converge t	rting th	e date d	column i					
	#thi #dai df3. df3 #Las df3.	s incon ly valu set_ind = df3.r t 10 va tail(10 D 2022-01 2022-01	istendes. ex('Daesamp. lues.) ate Ti -22 -23 -24	mmins 147.5	rting th	e date d	column i					
	#thi #dai df3. df3 #Las df3.	s incon ly valu set_ind = df3.r t 10 va tail(10 D 2022-01 2022-01	istendes. ex('Daesamp. lues.) ate Ti -22 -23 -24 -25	mmins 147.5 156.3	rting th	e date d	column i					
	#thi #dai df3. df3 #Las df3.	s incon ly valu set_ind = df3.r t 10 va tail(10 D 2022-01 2022-01 2022-01	istendes. ex('Daesamp. lues.) ate Ti -22 -23 -24 -25 -26	mmins 147.5 156.3 156.3	rting th	e date d	column i					
	#thi #dai df3. df3 #Las df3. 11707 11708 11710 11711 11712	s incon ly valu set_ind = df3.r t 10 va tail(10 D 2022-01 2022-01 2022-01 2022-01	istendes. ex('Daesamp. lues.) ate Ti -22 -23 -24 -25 -26 -27	mmins 147.5 156.3 156.3	rting th	e date d	column i					
	#thi #dai df3. df3 #Las df3. 11707 11708 11710 11711 11712 11713	s incon ly valu set_ind = df3.r t 10 va tail(10 D 2022-01 2022-01 2022-01 2022-01 2022-01	istendes. ex('Daesamp. lues.) ate Ti -22 -23 -24 -25 -26 -27 -28	mmins 147.5 156.3 156.3 156.3	rting th	e date d	column i					
•	#thi #dai df3. df3 #Las df3. 11707 11708 11710 11711 11712 11713 11714	s incon ly valu set_ind = df3.r t 10 va tail(10 D 2022-01 2022-01 2022-01 2022-01 2022-01 2022-01	istendes. ex('Daesamp. lues.) ate Ti -22 -23 -24 -25 -26 -27 -28 -29	mmins 147.5 156.3 156.3 156.3 156.3	rting th	e date d	column i					

11716 2022-01-31

156.2

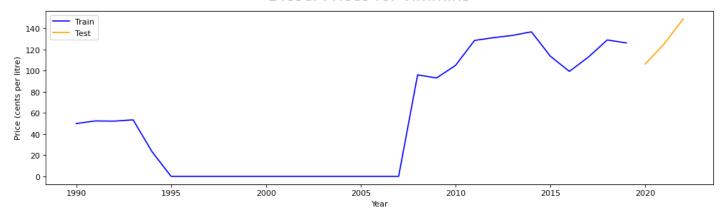
```
In [13]:
          df3.isnull().values.any()
         False
Out[13]:
In [14]:
          #Splitting the "Date" column into year, month and week to explore trends
          df3['Year']=df3['Date'].dt.year
          df3['Month'] = df3['Date'].dt.month
          df3['Week']=df3['Date'].dt.isocalendar().week
In [15]:
          df3.head(10)
                Date Timmins Year Month Week
Out[15]:
         0 1990-01-03
                         47.2 1990
         1 1990-01-04
                         47.2 1990
         2 1990-01-05
                         47.2 1990
                                              1
         3 1990-01-06
                         47.2 1990
                                              1
         4 1990-01-07
                         47.2 1990
                                              1
         5 1990-01-08
                         47.2 1990
                                              2
         6 1990-01-09
                         47.2 1990
                                              2
         7 1990-01-10
                                              2
                         47.4 1990
         8 1990-01-11
                         47.4 1990
                                              2
         9 1990-01-12
                         47.4 1990
                                              2
In [16]:
          #Splitting the dataset in Train and Test
          #Train from Year 1990 to Year 2019
          #Test from Year 2020
          train = df3[(df3['Date'] > '1990-01-01') & (df3['Date'] <= '2019-12-31')]
          test = df3[df3['Date'] >= '2020-01-01']
```

Yearly Price Visualization on Train and Test Dataset

#Checking for null values

```
In [17]:
         import matplotlib.pyplot as plt
         from matplotlib.pyplot import figure
         yearly train Price = train.groupby(['Year'])['Timmins'].mean()
         yearly test Price = test.groupby(['Year'])['Timmins'].mean()
         figure(figsize=(15, 4), dpi=80)
         plt.plot(yearly train Price, label='Train',c='blue')
         plt.plot(yearly test Price, label='Test',c='orange')
         plt.legend(loc='best')
         plt.suptitle('Diesel Prices for Timmins', fontsize=20)
         plt.xlabel('Year')
         plt.ylabel('Price (cents per litre)')
         plt.show()
```

Diesel Prices for Timmins



DataPrep for Time Series

```
In [18]:
    train.index = pd.DatetimeIndex(train['Date'])
    #Changing the frequency of the index to Daily
    train.index = train.asfreq('d').index

test.index = pd.DatetimeIndex(test['Date'])
#Changing the frequency of the index to Daily
    test.index = test.asfreq('d').index
```

Train and Time Series Dataset

```
In [19]: train_time_series = pd.DataFrame()
    train_time_series.index = train.index
    train_time_series.insert(0,"Timmins Diesel Price Train",train['Timmins'],True)

    test_time_series = pd.DataFrame()
    test_time_series.index = test.index
    test_time_series.insert(0,"Timmins Diesel Price Test",test['Timmins'],True)

    print(train_time_series.tail())
    print(test_time_series.tail())
```

```
Timmins Diesel Price Train
Date
                                   129.9
2019-12-27
2019-12-28
                                   129.9
2019-12-29
                                   129.9
2019-12-30
                                   129.9
2019-12-31
                                   129.9
            Timmins Diesel Price Test
Date
2022-01-27
                                  156.3
2022-01-28
                                  156.3
                                  156.3
2022-01-29
2022-01-30
                                  156.3
2022-01-31
                                  156.2
```

ARIMA Model

To predict time series with ARIMA, we need to set the values of three parameters (p,d,q):

p: The order of the auto-regressive (AR) model (i.e., the number of lag observations)

d: The degree of differencing.

q: The order of the moving average (MA) model.

Checking if data is stationary - We can see that it is not based on the P-value - Augmented Dickey Fuller Test

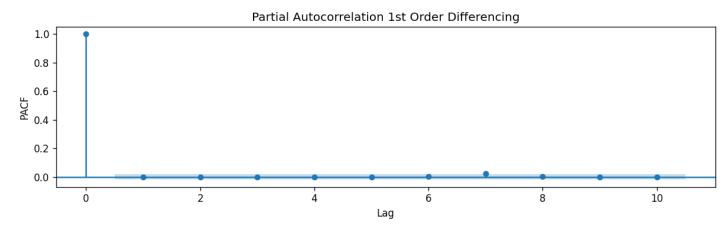
Taking First difference - P value is < 0.05. We can stop at the First Difference; d = 1

```
In [21]: train_time_series_stationary1 = train_time_series.diff().dropna()
    results1 = adfuller(train_time_series_stationary1['Timmins Diesel Price Train'])
    print('ADF Statistic: ',results1[0])
    print('p-value: ',results1[1])
    print('Critical Values', results1[4])

ADF Statistic: -104.65013792293605
    p-value: 0.0
    Critical Values {'1%': -3.4309471727572607, '5%': -2.861803917364605, '10%': -2.5669104780 8449}
```

The Order of Autoregressive Term p; p = 0

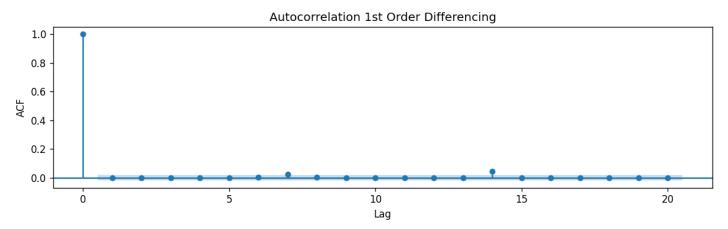
```
In [22]:
    plt.rcParams.update({'figure.figsize':(12,3), 'figure.dpi':120})
    from statsmodels.graphics.tsaplots import plot_pacf
    plot_pacf(train_time_series_stationary1, lags=10, title="Partial Autocorrelation 1st Order plt.xlabel('Lag')
    plt.ylabel('PACF')
    plt.show()
```



The order of the Moving Average term q; q = 0

In [23]:

```
plt.rcParams.update({'figure.figsize':(12,3), 'figure.dpi':120})
from statsmodels.graphics.tsaplots import plot_acf
plot_acf(train_time_series_stationary1, lags=20, title="Autocorrelation 1st Order Differe
plt.xlabel('Lag')
plt.ylabel('ACF')
plt.show()
```



Check the above p,d,q parameters with Auto Arima - Best Model ARIMA (0,1,0)

```
In [24]:
         import pmdarima as pm
         from pmdarima.model selection import train test split
         import numpy as np
         model1 = pm.auto arima(train time series, trace=True, error action='ignore', suppress warr
         model1.fit(train time series)
         forecast1 = model1.predict(n periods=len(test time series))
         forecast1 = pd.DataFrame(forecast1, index = test time series.index, columns=['Prediction'])
        Performing stepwise search to minimize aic
         ARIMA(2,1,2)(0,0,0)[0] intercept : AIC=38225.382, Time=1.52 sec
         ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=38217.382, Time=0.23 sec
         ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=38219.382, Time=0.68 sec
         ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=38219.382, Time=0.81 sec
                                            : AIC=38215.708, Time=0.13 sec
         ARIMA(0,1,0)(0,0,0)[0]
         ARIMA(1,1,1)(0,0,0)[0] intercept : AIC=38221.382, Time=1.00 sec
        Best model: ARIMA(0,1,0)(0,0,0)[0]
        Total fit time: 4.407 seconds
```

Model Summary

```
import statsmodels.api as sm
model = sm.tsa.arima.ARIMA(train_time_series, order=(0,1,0))
model_result = model.fit()
print(model_result.summary())
```

SARIMAX Results

```
______
Dep. Variable:
              Timmins Diesel Price Train No. Observations:
                                                              10955
                                                          -19106.854
Model:
                       ARIMA(0, 1, 0) Log Likelihood
Date:
                      Sun, 06 Mar 2022
                                   AIC
                                                           38215.708
Time:
                            12:32:16 BIC
                                                           38223.009
                          01-03-1990
                                                           38218.168
Sample:
                                   HQIC
                         - 12-31-2019
Covariance Type:
                                opg
```

=========	coef	std err	z	P> z	[0.025	0.975]
sigma2	1.9169	0.000	4326.192	0.000	1.916	1.918
Ljung-Box (L1 Prob(Q): Heteroskedast Prob(H) (two-	icity (H):		0.00 1.00 0.44 0.00	Jarque-Bera Prob(JB): Skew: Kurtosis:	(JB):	21310356346.66 0.00 66.55 6834.75
==========						

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

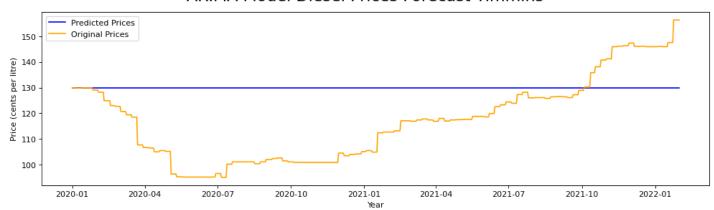
Model Prediction

```
In [26]: import warnings
    warnings.filterwarnings('ignore')
    ARIMA_Predict = model_result.predict(start='1/1/2020', end='1/31/2022')
    ARIMA_Predict_df = pd.DataFrame(ARIMA_Predict)
In [27]: ARIMA_Predict_df.tail()
```



```
In [28]:
    figure(figsize=(15, 4), dpi=80)
    plt.plot(ARIMA_Predict_df, label='Predicted Prices',c='blue')
    plt.plot(test_time_series, label='Original Prices',c='orange')
    plt.legend(loc='best')
    plt.suptitle('ARIMA Model Diesel Prices Forecast Timmins', fontsize=20)
    plt.xlabel('Year')
    plt.ylabel('Price (cents per litre)')
    plt.show()
```

ARIMA Model Diesel Prices Forecast Timmins



Evaluation of the Model

Mean Absolute Error (MAE) ARIMA

```
In [29]:
    from sklearn.metrics import mean_absolute_error
    maeARIMA=mean_absolute_error(test_time_series['Timmins Diesel Price Test'], ARIMA_Predict)
    print('Mean Absolute Error ARIMA = {}'.format(round(maeARIMA, 2)))

Mean Absolute Error ARIMA = 17.25
```

Mean squared error (MSE) ARIMA

```
In [30]:
    from sklearn.metrics import mean_squared_error
    mseARIMA=mean_squared_error(test_time_series['Timmins Diesel Price Test'], ARIMA_Predict)
    print('The Mean Squared Error ARIMA = {}'.format(round(mseARIMA, 2)))
The Mean Squared Error ARIMA = 412.19
```

Root mean squared error (RMSE) ARIMA

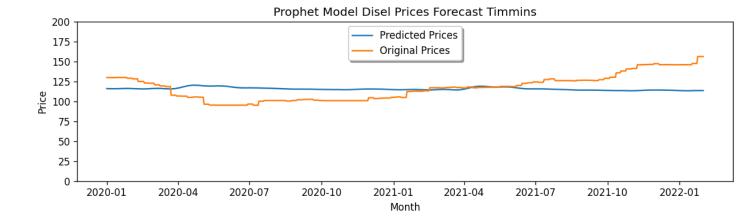
```
In [31]:
    from numpy import sqrt
    rmseARIMA = sqrt(mseARIMA)
    print('The Root Mean Squared Error ARIMA = {}'.format(round(rmseARIMA, 2)))
```

The Root Mean Squared Error ARIMA = 20.3

Prophet Model

plt.show()

```
In [32]:
         from fbprophet import Prophet
         d={'ds':train['Date'],'y':train['Timmins']}
         df pred=pd.DataFrame(data=d)
         model prophet = Prophet(daily seasonality=False)
         model prophet result = model prophet.fit(df pred)
In [33]:
         future = model prophet.make future dataframe(periods=765)
         forecast = model prophet.predict(future)
         forecast = forecast[(forecast['ds' ] \Rightarrow= '2020-01-01') & (forecast['ds' ] <= '2022-01-31')]
In [34]:
         fig, ax = plt.subplots()
         ax.plot(forecast['ds'], forecast['yhat'], label='Predicted Prices')
         ax.plot(test['Date'], test['Timmins'], label='Original Prices')
         plt.ylim([0,200])
         legend = ax.legend(loc='upper center', shadow=True)
         plt.title('Prophet Model Disel Prices Forecast Timmins')
         plt.xlabel('Month')
         plt.ylabel('Price')
```



Mean Absolute Error (MAE) Prophet

```
In [35]: maeProphet=mean_absolute_error(test['Timmins'],forecast['yhat'])
    print('Mean Absolute Error Prophet = {}'.format(round(maeProphet, 2)))

Mean Absolute Error Prophet = 13.87
```

Mean squared error (MSE) Prophet

Root mean squared error (RMSE) Prophet

```
In [37]:
    rmseProphet = sqrt(mseProphet)
    print('The Root Mean Squared Error Prophet = {}'.format(round(rmseProphet, 2)))
```

The Root Mean Squared Error Prophet = 16.86